
A Warehouse Floor at Scale

Three TurtleBots under Open-RMF Executing an
Epistemic Policy over Forty Metres of Warehouse

Haniel Ulises

Epistemic Robotics

September 8, 2026

The companion report executes an epistemic mission on a floor of fourteen metres by twenty-one with two aisles. This one asks what changes when the floor is forty-two by sixty-three, holds a hundred and fifty-three obstacles and nine people walking through it, and is driven by three robots rather than two. Nothing in the epistemic layer changes, which is the point; a great deal in the execution layer does, and most of this report is an account of what.

Contents

1	Scope, and what this report is for	2
2	The floor	2
2.1	Why an existing floor, and which	2
2.2	Deriving the roadmap	2
2.3	Aisles are found, not listed	3
3	The fleet	3
3.1	A TurtleBot3 that Open-RMF can drive	3
3.2	Pinned allocation	4
4	The mission	4
5	Defects found by the integration	4
6	What forty metres does not establish	6
7	Conclusion	6

1 Scope, and what this report is for

The claim under examination is narrow and worth stating before anything else. It is *not* that epistemic planning scales: the policy solved here has four nodes and would be solved as quickly over a floor plan drawn on a napkin. It is that the execution layer between such a policy and a real fleet holds together at a scale where the floor is no longer trivially traversable, and that the properties the policy is chosen for — in particular, that one agent must *not* come to know something — remain checkable when it does.

Two things follow from that framing. First, the interesting content of this report is the failure analysis of Section 5, not the mission trace of Section 4: the mission is three actions long and its trace could be summarised in a sentence. Second, the negative results matter as much as the positive ones, and Section 6 states plainly what forty metres of warehouse does not establish.

2 The floor

2.1 Why an existing floor, and which

Open-RMF routes over a navigation graph and a building map, and its demonstrations ship both. The largest maps in `rmf_demos` are the airport terminal, at 279×60 m with 92 named waypoints, and a campus at 319×246 m; neither is a warehouse. There is no large warehouse map published for Open-RMF at all.

The candidate warehouse assets on Gazebo Fuel were therefore measured directly, from their own collision meshes rather than from their catalogue descriptions. Table 1 records the result, and the discrepancy between what the descriptions promise and what the geometry provides is the substance of it.

Table 1: Warehouse assets on Gazebo Fuel, measured from their collision meshes.

Asset	Footprint	Collision	Consequence
OpenRobotics/Depot	29×15 m	a ground plane	Racks and walls are visuals only. A laser passes through them and an occupancy grid rasterised from collision geometry returns an empty floor.
OpenRobotics/Warehouse	30×50 m	shell and pillars	Genuinely large and genuinely empty: 1.2% occupied at robot height, all of it structural columns.
industrial-warehouse	as AWS small	complete	The AWS small warehouse recomposed from the same parts; no larger than the floor already in use.

Depot is the asset usually reached for — some fifty thousand downloads — and is the trap of the three. Its detail is entirely visual.

The floor finally used is `dynamic_logistics_warehouse`, which is AWS RoboMaker’s small warehouse tiled and furnished by hand: nine floor slabs, 153 obstacles drawn from fourteen AWS assets, five lamps, and nine scripted actors walking fixed circuits. Measured from its own model poses it spans 42.0×62.8 m, against 14.05×21.05 m for the floor of the companion report — an increase of roughly nine times in area.

2.2 Deriving the roadmap

The companion report’s floor is parametric and its navigation graph is written from a formula. This one is not: nothing about the placement of 153 hand-set obstacles is regular, so the graph is derived from the world instead.

1. Every model in the world is one of fourteen AWS assets. Their collision extents are measured once, from the collision meshes, and cached.
2. Three of the fourteen are structure rather than obstacle. Getting this wrong is not a detail: `GroundB` is a 14×21 m floor tile, and treating it as an obstacle fills the warehouse solid.
3. The remaining models are rasterised at 0.25 m and inflated by the robot’s radius plus a margin, giving 0.45 m of clearance.
4. A lattice is laid over the free floor, lanes are kept only where the straight segment between their ends is clear, and the roadmap is then thinned.

The thinning is where the engineering is. A node is removed whenever every pair of its neighbours can be joined directly, which preserves every route and costs one waypoint; junctions and aisle waypoints are kept. Table 2 shows why both halves of that — fine lattice, then thinning — are necessary.

Table 2: Lattice spacing against connectivity and tractability.

Spacing	Free nodes	Largest component	Outcome
2.00 m	295	147 (50 %)	floor splits into three bands
1.50 m	541	447 (83 %)	still disconnected
1.25 m	748	748 (100 %)	connected; planner intractable
1.00 m	1205	1205 (100 %)	connected; larger still
1.25 m, thinned	261	261 (100 %)	connected and routable

Six of the racks are eighteen metres long, so the floor is crossed by walls with few gaps. At two-metre spacing the lattice misses enough of those gaps to leave the warehouse in three disconnected north–south bands, and a largest-component filter then silently discards two thirds of the building. At 1.25 m every gap is found. The thinning recovers what the fine spacing costs: 261 waypoints, 35 of them named as aisles, and 1652 directed lanes.

2.3 Aisles are found, not listed

A waypoint is called an aisle when the floor runs out close on both sides along one axis and is open along the other. That is exactly the place a robot must enter before it can tell anything about what is inside, and so exactly what a sensing action wants to name. Thirty-five such waypoints survive a four-metre minimum separation, one per aisle rather than one per cell.

3 The fleet

3.1 A TurtleBot3 that Open-RMF can drive

The fleet is three TurtleBot3 Waffles. Assembling one that RMF can actually drive took two corrections, both of which are worth recording because neither is documented anywhere obvious.

The skeleton. RMF moves a simulated robot with the `slotcar` plugin, which is a kinematic driver rather than a controller. A Waffle built from scratch to what `slotcar` documents it requires — a `base_footprint`, a `left_wheel` and `right_wheel`, and the revolute joints `joint_tire_left` and `joint_tire_right` — loaded without complaint, initialised `slotcar`, published its state at 6 Hz, accepted paths from the fleet adapter, and did not move a centimetre. On the same floor and the same graph, `rmf_demos`’ `TinyRobot` drove sixty metres without trouble, which is what isolated the model as the fault rather than the world or the roadmap. The Waffle geometry is therefore mounted on `TinyRobot`’s drive skeleton unchanged.

The scale. `turtlebot3_gazebo` ships its meshes as STL in millimetres. Loaded at scale 1 a Waffle base is 271 metres across. The `.dae` files beside them in `turtlebot3_common` are placeholder cubes two metres on a side, which is a second and independent way of producing an enormous robot. Both were encountered.

3.2 Pinned allocation

Tasks are pinned to named robots through `robot_task_request` rather than dispatched for bidding. The reason is epistemic rather than practical: the planner assigns actions to named agents while it solves, the model records knowledge against those names, and if the dispatcher were free to substitute a robot the model would credit an agent with having sensed something it never saw. An agent bound to no robot is a hard error and the bridge refuses the action.

4 The mission

The domain is the survey of `eplansys`, unchanged. Three agents; a site that may be contaminated; a goal of three conjuncts:

$$\gamma = Kw_{scout} \text{ contaminated} \wedge Kw_{relay} \text{ contaminated} \wedge \neg Kw_{observer} \text{ contaminated}.$$

The third conjunct is what gives the domain its difficulty. It forbids the open channel, which is otherwise the cheapest way for the team to inform itself, and so forces the planner to choose a channel by its audience. The planner returns a policy of four nodes, branching.

What the larger floor changes is the geography the policy is executed over. The site is `aisle_20`, thirty metres from the robot's charger and reachable only by way of the service floor: an aisle cannot be seen into until a robot stands at its mouth, which is the physical counterpart of the sensing action the policy contains.

5 Defects found by the integration

Putting the chain in operation at this scale revealed eight defects that no part of it exhibited alone. They are given in full because, taken together, they are the substantive result of the exercise: six of the eight are consequences of composition rather than errors of programming, and none would have appeared in a unit test.

Table 3: Defects encountered, in the order they were met.

Symptom	Cause	Nature
No robot moves; the fleet manager is absent.	The system <code>fastapi</code> is 0.63, written against pydantic 1; a pydantic 2 in <code>~/local</code> shadows the packaged 1.8 for every interpreter on the machine and the import fails.	Environment
A robot leaves its charger, stops after 1.3 m, and RMF reports the task underway.	A pallet jack placed for appearance overlapped the charger’s own lane by 20 cm. The graph check knew about racks and pillars and not about props.	Composition
The warehouse is a third of its size.	The nine floor slabs do not abut; a 12 cm seam reads as absent floor and splits the map into three components, of which a largest-component filter keeps one.	Geometric
No path is returned for a 60 m errand in 70 s.	693 waypoints and 2200 lanes. The fine lattice needed for connectivity is not the roadmap that should be routed on.	Scale
Robots report a pose, accept paths and do not move.	Spawned at $z = 0$ into floor slabs 12 cm thick with no ground plane beneath; each came to rest at $z = -0.067$ with its wheels inside the floor.	Geometric
A purpose-built TurtleBot3 loads, initialises slotcar, publishes state and does not move.	slotcar requires more of a model than it documents.	Undocumented interface
The robot is 271 m across; or, separately, a 2 m cube.	<code>turtlebot3_gazebo</code> meshes are STL in millimetres, and the <code>.dae</code> files beside them are placeholders.	Units
<code>goto_site</code> fails with <i>RMF task timed out</i> .	The 180 s limit is the small floor’s figure. Forty metres at half a metre per second is not 180 s.	Scale

Three of these deserve comment.

The pallet jack. The graph check was written to refuse lanes that cross obstacles, and it did refuse three layouts that looked correct on paper. It passed this one because its obstacle set was racks and pillars: the prop had been placed in the world generator, for appearance, and the checker never saw it. The symptom was severe out of all proportion to the cause — a robot pressed motionless into a pallet jack while RMF reported its task underway and nothing anywhere reported a collision. The correction moves the props into the same definition the checker reads, and the check that missed it now rejects the original placement.

The seam. A twelve-centimetre gap between two floor tiles is invisible in any rendering and fatal to a rasteriser, which cannot distinguish a seam from a hole. What makes it worth recording is the interaction with the largest-component filter: the filter is correct in isolation — an isolated pocket of floor is a waypoint RMF will accept and never route to — and together they discard two thirds of a warehouse without a diagnostic.

The lattice. Connectivity and tractability pull in opposite directions, and the resolution that satisfies one fails the other. Building finely and thinning afterwards is not the same as building coarsely, and Table 2 is the measurement that says so.

6 What forty metres does not establish

Not a planning result. The policy has four nodes. Nothing here bears on the complexity of epistemic planning, and the floor could be enlarged indefinitely without touching it.

Not a study of contention. Three robots share one floor and three tasks were accepted and executed concurrently without a process failure, but no measurement was made of how the traffic schedule mediates the shared corridors under sustained multi-robot traffic. That is the obvious next experiment and it has not been run.

Not a dynamic-obstacle result. The world's nine actors walk fixed circuits and are not in the navigation graph. The fleet does not avoid them and was never asked to; they are scenery, and a robot that met one would drive into it.

Not confidentiality. The private channel is realised by addressing: an utterance is published to a topic whose subscribers are the audience the action declares. Any node may subscribe to any topic. What is established is that the team used a channel the excluded agent was not addressed on, not that it could not have listened.

7 Conclusion

The execution layer holds at this scale, and the eight defects of Section 5 are the price of finding that out. Six of them are properties of the composition rather than of any component, which is the recurring lesson of this line of work: a system of autonomous nodes, plugins loaded at run time and streams advancing at different rates exhibits a family of faults that integration alone can reveal.

The epistemic content is unchanged from the small floor, and deliberately so. What the larger floor buys is that the geography is no longer trivial — the site is thirty metres away and down an aisle — so that the claim the domain exists to support, that an agent can be kept from coming to know something, is made against a floor where getting anywhere at all is work.